

A Fast Elitist Non-dominated Sorting Genetic Algorithm for Multi-objective Optimization: NSGA-II

Kalyanmoy Deb, Samir Agrawal, Amrit Pratap, and T Meyarivan

Kanpur Genetic Algorithms Laboratory (KanGAL)
Indian Institute of Technology Kanpur
Kanpur, PIN 208 016, India
{deb,samira,apratap,mary}@iitk.ac.in
<http://www.iitk.ac.in/kangal>

Abstract. Multi-objective evolutionary algorithms which use non-dominated sorting and sharing have been mainly criticized for their (i) $O(MN^3)$ computational complexity (where M is the number of objectives and N is the population size), (ii) non-elitism approach, and (iii) the need for specifying a sharing parameter. In this paper, we suggest a non-dominated sorting based multi-objective evolutionary algorithm (we called it the Non-dominated Sorting GA-II or NSGA-II) which alleviates all the above three difficulties. Specifically, a fast non-dominated sorting approach with $O(MN^2)$ computational complexity is presented. Second, a selection operator is presented which creates a mating pool by combining the parent and child populations and selecting the best (with respect to fitness and spread) N solutions. Simulation results on five difficult test problems show that the proposed NSGA-II, in most problems, is able to find much better spread of solutions and better convergence near the true Pareto-optimal front compared to PAES and SPEA—two other elitist multi-objective EAs which pay special attention towards creating a diverse Pareto-optimal front. Because of NSGA-II's low computational requirements, elitist approach, and parameter-less sharing approach, NSGA-II should find increasing applications in the years to come.

1 Introduction

Over the past decade, a number of multi-objective evolutionary algorithms (MOEAs) have been suggested [8, 3, 5, 11]. The primary reason for this is their ability to find multiple Pareto-optimal solutions in one single run. Since the principal reason why a problem has a multi-objective formulation is because it is not possible to have a single solution which simultaneously optimizes all objectives, an algorithm that gives a large number of alternative solutions lying on or near the Pareto-optimal front is of great practical value.

The Non-dominated Sorting Genetic Algorithm (NSGA) proposed in Srinivas and Deb [8] was one of the first such evolutionary algorithms. Over the years, the main criticism of the NSGA approach have been as follows:

High computational complexity of non-dominated sorting: The non-dominated sorting algorithm in use until now is $O(MN^3)$ which in case of large population sizes is very expensive, especially since the population needs to be sorted in every generation.

Lack of elitism: Recent results [10, 7] show clearly that elitism can speed up the performance of the GA significantly, also it helps to prevent the loss of good solutions once they have been found.

Need for specifying the sharing parameter σ_{share} : Traditional mechanisms of insuring diversity in a population so as to get a wide variety of equivalent solutions have relied heavily on the concept of sharing. The main problem with sharing is that it requires the specification of a sharing parameter (σ_{share}). Though there has been some work on dynamic sizing of the sharing parameter [4], a parameterless diversity preservation mechanism is desirable.

In this paper, we address all of these issues and propose a much improved version of NSGA which we call NSGA-II. From the simulation results on a number of difficult test problems, we find that NSGA-II is, in general, better than PAES and SPEA—two other elitist multi-objective evolutionary algorithm—in terms of converging near the Pareto-optimal front and maintaining diversity among obtained solutions. These results encourage the application of NSGA-II to more complex and real-world multi-objective optimization problems.

2 Elitist Multi-Objective Evolutionary Algorithms

In the study of Zitzler, Deb, and Thiele [10], it was clearly shown that elitism helps in achieving better convergence in MOEAs. Among the existing elitist MOEAs, Zitzler and Thiele's [11] strength Pareto EA (SPEA), Knowles and Corne's Pareto-archived evolution strategy (PAES) [6], and Rudolph's [7] elitist GA are well known.

Zitzler and Thiele [11] suggested an elitist multi-criterion EA with the concept of non-domination in their strength Pareto EA (SPEA). They suggested maintaining an external population at every generation storing all non-dominated solutions discovered so far beginning from the initial population. This external population participates in genetic operations. At each generation, a combined population with the external and the current population is first constructed. All non-dominated solutions in the combined population are assigned a fitness based on the number of solutions they dominate and dominated solutions are assigned fitness worse than the worst fitness of any non-dominated solution. This assignment of fitness makes sure that the search is directed towards the non-dominated solutions. A deterministic clustering technique is used to ensure diversity among non-dominated solutions. Although the implementation suggested in [11] is $O(MN^3)$, with proper book-keeping the complexity of SPEA can be reduced to $O(MN^2)$.

Knowles and Corne [6] suggested a simple MOEA using an evolution strategy (ES). In their Pareto-archived ES (PAES) with one parent and one child, the child is compared with respect to the parent. If the child dominates the parent, the child is accepted as the next parent and the iteration continues. On the other hand, if the parent dominates the child, the child is discarded and a new mutated solution (a new child) is found. However, if the child and the parent do not dominate each other, the choice between the child and the parent is made by comparing them with an archive of best solutions found so far. The child is compared with the archive to check if it dominates any member of the archive. If yes, the child is accepted as the new parent and the dominated solution is

eliminated from the archive. If the child does not dominate any member of the archive, both parent and child are checked for their *nearness* with the solutions of the archive. If the child resides in a least crowded region in the parameter space among the members of the archive, it is accepted as a parent and a copy of added to the archive. Authors have calculated the worst case complexity of PAES for N evaluations as $O(aMN)$, where a is the archive length. Since the archive size is usually chosen proportional to the population size N , the overall complexity of the algorithm is $O(MN^2)$.

Rudolph [7] suggested, but did not simulate, a simple elitist multi-objective EA based on a systematic comparison of individuals from parent and offspring populations. The non-dominated solutions of the offspring population are compared with that of parent solutions to form an overall non-dominated set of solutions, which becomes the parent population of the next iteration. If the size of this set is not greater than the desired population size, other individuals from the offspring population are included. With this strategy, he has been able to prove the convergence of this algorithm to the Pareto-optimal front. Although this is an important achievement in its own right, the algorithm lacks motivation for the second task of maintaining diversity of Pareto-optimal solutions.

3 Elitist Non-dominated Sorting Genetic Algorithm (NSGA-II)

The non-dominated sorting GA (NSGA) proposed by Srinivas and Deb in 1994 has been subjected to a number of criticism, as mentioned earlier. In this section, we suggest NSGA-II, which alleviate all these difficulties. We begin by presenting a number of different modules that form part of NSGA-II.

3.1 A fast non-dominated sorting approach

In order to sort a population of size N according to the level of non-domination, each solution must be compared with every other solution in the population to find if it is dominated. This requires $O(MN)$ comparisons for each solution, where M is the number of objectives. When this process is continued to find the members of the first non-dominated class for all population members, the total complexity is $O(MN^2)$. At this stage, all individuals in the first non-dominated front are found. In order to find the individuals in the next front, the solutions of the first front are temporarily discounted and the above procedure is performed again. The procedure is repeated to find the subsequent fronts. As can be seen the worst case (when there exists only one solution in each front) complexity of this algorithm without any book-keeping is $O(MN^3)$. In the following we describe a fast non-dominated sorting approach which will require at most $O(MN^2)$ computations.

This approach is similar in principle to above approach, except that a better book-keeping strategy is performed to make it a faster algorithm. In this approach, every solution from the population is checked with a partially filled population for domination. To start with, the first solution from the population is kept in a set P' . Thereafter, each solution p (the second solution onwards) is compared with all members of the set P' one by one. If the solution p dominates any member q of P' , then solution q

is removed from P' . This way non-members of the non-dominated front get deleted from P' . Otherwise, if solution p is dominated by any member of P' , the solution p is ignored. If solution p is not dominated by any member of P' , it is entered in P' . This is how the set P' grows with non-dominated solutions. When all solutions of the population are checked, the remaining members of P' constitute the non-dominated set.

`fast-nondominated-sort(P)`

$P' = \{1\}$	include first member in P'
for each $p \in P \wedge p \notin P'$	take one solution at a time
$P' = P' \cup \{p\}$	include p in P' temporarily
for each $q \in P' \wedge q \neq p$	compare p with other members of P'
if $p \prec q$, then $P' = P' \setminus \{q\}$	if p dominates a member of P' , delete it
else if $q \prec p$, then $P' = P' \setminus \{p\}$	if p is dominated by other members of P' ,
	do not include p in P'

To find other fronts, the members of P' will be discounted and the above procedure is repeated.

Here, we observe that the second element of the population is compared with only one solution P' , the third solution with at most two solutions of P' , and so on. This requires a maximum of $O(N^2)$ domination checks. Since each domination check requires M function value comparisons, the maximum complexity of this approach is also $O(MN^2)$.

3.2 Density estimation

To get an estimate of the density of solutions surrounding a particular point in the population we take the average distance of the two points on either side of this point along each of the objectives. This quantity $i_{distance}$ serves as an estimate of the size of the largest cuboid enclosing the point i without including any other point in the population (we call this the *crowding distance*). In Figure 1, the crowding distance of the i -th solution in its front (marked with solid circles) is the average side-length of the cuboid (shown with a dashed box). The following algorithm is used to calculate the crowding distance of each point in the set \mathcal{I} :

`crowding-distance-assignment(\mathcal{I})`

$l = \mathcal{I} $	number of solutions in \mathcal{I}
for each i , set $\mathcal{I}[i]_{distance} = 0$	initialize distance
for each objective m	
$\mathcal{I} = \text{sort}(\mathcal{I}, m)$	sort using each objective value
$\mathcal{I}[1]_{distance} = \mathcal{I}[l]_{distance} = \infty$	so that boundary points are always selected
for $i = 2$ to $(l - 1)$	for all other points
$\mathcal{I}[i]_{distance} = \mathcal{I}[i]_{distance} + (\mathcal{I}[i + 1].m - \mathcal{I}[i - 1].m)$	

Here $\mathcal{I}[i].m$ refers to the m -th objective function value of the i -th individual in the set \mathcal{I} . The complexity of this procedure is governed by the sorting algorithm. In the worst case (when all solutions are in one front), the sorting requires $O(mN \log N)$ computations.




Fig. 1. The crowding distance calculation is shown

3.3 Crowded comparison operator

The crowded comparison operator (\prec_n) guides the selection process at the various stages of the algorithm towards a uniformly spread-out Pareto-optimal front. Let us assume that every individual i in the population has two attributes.

1. Non-domination rank (i_{rank})
2. Local crowding distance ($i_{distance}$)

We now define a partial order \prec_n as :

$$i \prec_n j \quad \text{if } (i_{rank} < j_{rank}) \text{ or } ((i_{rank} = j_{rank}) \text{ and } (i_{distance} > j_{distance}))$$

That is, between two solutions with differing non-domination ranks we prefer the point with the lower rank. Otherwise, if both the points belong to the same front then we prefer the point which is located in a region with lesser number of points (the size of the cuboid inclosing it is larger).

3.4 The main loop

Initially, a random parent population P_0 is created. The population is sorted based on the non-domination. Each solution is assigned a fitness equal to its non-domination level (1 is the best level). Thus, minimization of fitness is assumed. Binary tournament selection, recombination, and mutation operators are used to create a child population Q_0 of size N . From the first generation onward, the procedure is different. The elitism procedure for $t \geq 1$ and for a particular generation is shown in the following:

$R_t = P_t \cup Q_t$	combine parent and children population
$\mathcal{F} = \text{fast-nondominated-sort}(R_t)$	$\mathcal{F} = (\mathcal{F}_1, \mathcal{F}_2, \dots)$, all non-dominated fronts of R_t
$P_{t+1} = \emptyset$	
until $ P_{t+1} < N$	till the parent population is filled
crowding-distance-assignment(\mathcal{F}_i)	calculate crowding distance in \mathcal{F}_i
$P_{t+1} = P_{t+1} \cup \mathcal{F}_i$	include i -th non-dominated front in the parent pop
Sort(P_{t+1}, \prec_n)	sort in descending order using \prec_n

$P_{t+1} = P_{t+1}[0 : N]$	choose the first N elements of P_{t+1}
$Q_{t+1} = \text{make-new-pop}(P_{t+1})$	use selection, crossover and mutation to create
$t = t + 1$	a new population Q_{t+1}

First, a combined population $R_t = P_t \cup Q_t$ is formed. The population R_t will be of size $2N$. Then, the population R_t is sorted according to non-domination. The new parent population P_{t+1} is formed by adding solutions from the first front till the size exceeds N . Thereafter, the solutions of the last accepted front are sorted according to \prec_n and a total of N solutions are picked. This is how we construct the population P_{t+1} . This population of size N is now used for selection, crossover and mutation to create a new population Q_{t+1} of size N . It is important to note that we use a binary tournament selection operator but the selection criterion is now based on the niched comparison operator \prec_n .

Let us now look at the complexity of one iteration of the entire algorithm. The basic operations being performed and the worst case complexities associated with are as follows:

1. Non-dominated sort is $O(MN^2)$,
2. Crowding distance assignment is $O(MN \log N)$, and
3. Sort on \prec_n is $O(2N \log(2N))$.

As can be seen, the overall complexity of the above algorithm is $O(MN^2)$.

The diversity among non-dominated solutions is introduced by using the crowding comparison procedure which is used in the tournament selection and during the population reduction phase. Since solutions compete with their crowding distance (a measure of density of solutions in the neighborhood), no extra niching parameter (such as σ_{share} needed in the NSGA) is required here. Although the crowding distance is calculated in the objective function space, it can also be implemented in the parameter space, if so desired [1].

4 Results

We compare NSGA-II with PAES on five test problems (minimization of both objectives) [9, 10]:

$$\text{MOP2: } \begin{cases} f_1(x) = 1 - \exp \left(- \sum_{i=1}^3 \left(x_i - \frac{1}{\sqrt{3}} \right)^2 \right) \\ f_2(x) = 1 - \exp \left(- \sum_{i=1}^3 \left(x_i + \frac{1}{\sqrt{3}} \right)^2 \right) \end{cases} \quad -4 \leq x_1, x_2, x_3 \leq 4 \quad (1)$$

$$\text{MOP3: } \begin{cases} f_1(x) = [1 + (A_1 - B_1)^2 + (A_2 - B_2)^2] \\ f_2(x) = [(x + 3)^2 + (y + 1)^2] \end{cases} \quad (2)$$

$$\begin{aligned} \text{where} \quad & A_1 = 0.5 \sin 1 - 2 \cos 1 + \sin 2 - 1.5 \cos 2 \\ & A_2 = 1.5 \sin 1 - \cos 1 + 2 \sin 2 - 0.5 \cos 2 \\ & B_1 = 0.5 \sin x - 2 \cos x + \sin y - 1.5 \cos y \\ & B_2 = 1.5 \sin x - \cos x + 2 \sin y - 0.5 \cos y \end{aligned}$$

$$\text{MOP4: } \begin{cases} f_1(\mathbf{x}) = \sum_{i=1}^{n-1} \left(-10 \exp \left(-0.2 \sqrt{x_i^2 + x_{i+1}^2} \right) \right) & -5 \leq x_1, x_2, x_3 \leq 5 \\ f_2(\mathbf{x}) = \sum_{i=1}^n (|x_i|^{0.8} + 5 \sin(x_i)^3) \end{cases} \quad (3)$$

$$\text{TC4: } \begin{cases} f_1(\mathbf{x}) = x_1 & 0 \leq x_1 \leq 1 \\ f_2(\mathbf{x}) = g \left(1 - \sqrt{\frac{x_1}{g}} \right) & -5 \leq x_2, \dots, x_{10} \leq 5 \end{cases} \quad (4)$$

$$\text{where } g(\mathbf{x}) = 91 + \sum_{i=2}^{10} (x_i^2 - 10 \cos(4\pi x_i))$$

$$\text{TC6: } \begin{cases} f_1(\mathbf{x}) = 1 - \exp(-4x_1) \sin^6(6\pi x_1) & 0 \leq x_i \leq 1 \quad i = 1, \dots, 10 \\ f_2(\mathbf{x}) = g \left(1 - (f_1/g)^2 \right) \end{cases} \quad (5)$$

$$\text{where } g(\mathbf{x}) = 1 + 9 \left(\sum_{i=2}^{10} x_i/9 \right)^{0.25}$$

Since the diversity among optimized solutions is an important matter in multi-objective optimization, we devise two measures—one based on the consecutive distances among the solutions of the best non-dominated front in the final population and the other based on the average distance of solutions from the known global Pareto-optimal front. The obtained set of the first non-dominated solutions are compared with a uniform distribution and the deviation is computed as follows:

$$\Delta = \sum_{i=1}^{|\mathcal{F}_1|} \frac{|d_i - \bar{d}|}{|\mathcal{F}_1|}. \quad (6)$$

In order to ensure that this calculation takes into account the spread of solutions in the entire region of the true front, we include the boundary solutions¹ in the non-dominated front \mathcal{F}_1 . For discrete Pareto-optimal fronts, we calculate a weighted average of the above metric for each of the discrete regions. In the above equation, d_i is the Euclidean distance between two consecutive solutions in the first non-dominated front of the final population in the objective function space. The parameter \bar{d} is the average of these distances.

The second metric γ measures the convergence property of an algorithm. From each solution in the non-dominated front, its perpendicular distance to the global Pareto-optimal front is calculated by approximating the Pareto-optimal front as a combination of 500 piece-wise linear segments. The average of these perpendicular distances is measured.

For all test problems and with NSGA-II, we use a population of size 100, a crossover probability of 0.8, a mutation probability of $1/n$ (where n is the number of variables). We run NSGA-II for 250 generations. The variables are treated as real numbers and the simulated binary crossover (SBX-20) [2] and the real-parameter mutation operator (with distribution index of 500) are used. For the (1+1)-PAES, we have used an archive size of 100 and depth of 4 [6]. For SPEA, we have used $N = 80$ and an external population of size 20. A crossover probability of 0.8 is used. A mutation probability of

¹ Boundary solutions are not considered for TC4 and TC6. This is because in these problems the none of the algorithms has converged to the global Pareto-optimal front.

0.01 is used² are different. In order to make the comparisons fair, we have used 25,000 iterations in PAES, so that total number of function evaluations in NSGA-II, PAES, and SPEA are the same.

Table 1 shows the deviation from an ideal (uniform) spread ($\bar{\Delta}$) and its variance in 10 independent runs. We show two columns for each test problem. The first column presents the $\bar{\Delta}$ value of 10 runs and the second column shows its variance. It is clear from the table that in most test problems NSGA-II has found much smaller $\bar{\Delta}$, meaning that NSGA-II is able to find a distribution of solutions closer to a uniform distribution along the non-dominated front. The variance columns suggest that the obtained Δ values are consistent in all 10 runs. Table 2 shows the average of convergence measure γ

Table 1. Comparison of mean and variance of deviation measure Δ obtained using NSGA-II, PAES, and SPEA

Algorithm	MOP2		MOP3		MOP4		TC4		TC6	
NSGA-II	0.361	0.00068	0.445	0.00043	0.387	0.00164	0.383	0.00099	0.365	0.01613
PAES	1.609	0.00671	1.341	0.00495	1.087	0.00687	1.563	0.05723	1.195	0.05151
SPEA	0.740	0.00748	0.880	0.00508	0.733	0.00175	0.167	0.00000	0.804	0.01142

and its standard deviation of 10 runs. It is evident that NSGA-II can come closer to the actual Pareto-optimal front in all problems compared to other two algorithms (except in TC4, where PAES is the best).

Table 2. Comparison of distance γ from the true Pareto-optimal front and its standard deviation obtained using NSGA-II, PAES and SPEA

Algorithm	MOP2		MOP3		MOP4		TC4		TC6	
NSGA-II	0.0019	0.00000	0.0151	0.00000	0.0250	0.00004	4.5128	4.55386	0.0611	0.00056
PAES	0.1704	0.00002	11.8315	12.13053	0.1046	0.03054	0.5846	0.53599	0.1999	0.00122
SPEA	0.1257	0.00004	0.0378	0.00009	0.0456	0.00005	7.3403	6.57252	0.2211	0.00045

In order to have a better understanding of how these algorithms are able to spread solutions over the non-dominated front, we present the entire non-dominated front found by NSGA-II and PAES in two of the above five test problems. Figures 2 and 3 show that NSGA-II is able to find a much better distribution than PAES on MOP4. In TC4, converging to the global Pareto-optimal front is a difficult task. PAES's grid assignment scheme does well in coming closer to the global Pareto-optimal front. With NSGA-II,

² Mutation probability used in PAES and SPEA is different from NSGA-II, since in NSGA-II a real-parameter mutation operator is used and in PAES and SPEA a bit-wise mutation operator is used.




Fig. 2. Non-dominated solutions obtained using NSGA-II on MOP4




Fig. 3. Non-dominated solutions obtained using PAES on MOP4

we find a front with $g = 3.5$ in one out of five different runs. Figure 4 shows the non-dominated solutions obtained using NSGA-II, PAES, and SPEA for TC6. It is clear that the NSGA-II is able to better distribute its population along the obtained front than PAES.

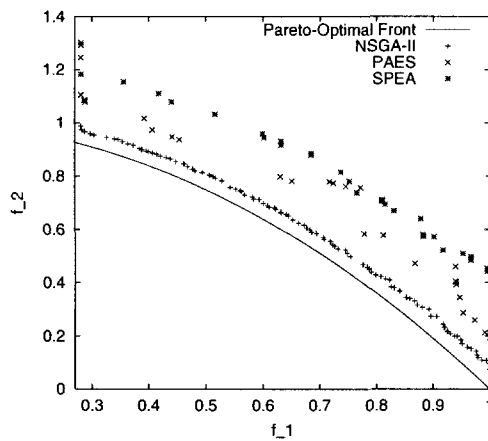


Fig. 4. Obtained non-dominated solutions with NSGA-II and PAES on TC6

5 Conclusions

In this paper, we have proposed a computationally fast elitist multi-objective evolutionary algorithm based on non-dominated sorting approach. On five difficult test problems

borrowed from the literature, it has been found that the proposed NSGA-II outperforms PAES and SPEA—two other popular multi-objective EAs with the explicit goals of preserving spread on the non-dominated front. With the properties of a fast non-dominated sorting procedure, an elitist strategy, and a parameterless approach, NSGA-II should find increasing attention and applications in the near future.

Acknowledgements

Authors acknowledge the support provided by All India Council for Technical Education, India during the course of this study.

References

1. Deb, K. (1999) Multi-objective genetic algorithms: Problem difficulties and construction of test Functions. *Evolutionary Computation*, 7(3), 205–230.
2. Deb, K. and Agrawal, R. B. (1995) Simulated binary crossover for continuous search space. *Complex Systems*, 9 115–148.
3. Fonseca, C. M. and Fleming, P. J. (1993) Genetic algorithms for multi-objective optimization: Formulation, discussion and generalization. In Forrest, S., editor, *Proceedings of the Fifth International Conference on Genetic Algorithms*, pages 416–423, Morgan Kauffman, San Mateo, California.
4. Fonseca, C. M. and Fleming, P. J. (1998). Multiobjective optimization and multiple constraint handling with evolutionary algorithms—Part II: Application example. *IEEE Transactions on Systems, Man, and Cybernetics: Part A: Systems and Humans*. 38–47.
5. Horn, J. and Nafpliotis, N., and Goldberg, D. E. (1994) A niched Pareto genetic algorithm for multi-objective optimization. In Michalewicz, Z., editor, *Proceedings of the First IEEE Conference on Evolutionary Computation*, pages 82–87, IEEE Service Center, Piscataway, New Jersey.
6. Knowles, J. and Corne, D. (1999) The Pareto archived evolution strategy: A new baseline algorithm for multiobjective optimisation. *Proceedings of the 1999 Congress on Evolutionary Computation*, Piscataway: New Jersey: IEEE Service Center, 98–105.
7. Rudolph, G. (1999) Evolutionary search under partially ordered sets. Technical Report No. CI-67/99, Dortmund: Department of Computer Science/LS11, University of Dortmund, Germany.
8. Srinivas, N. and Deb, K. (1995) Multi-Objective function optimization using non-dominated sorting genetic algorithms, *Evolutionary Computation*, 2(3):221–248.
9. van Veldhuizen, D. and Lamont, G. B. (1998). Multiobjective evolutionary algorithm research: A history and analysis. *Report Number TR-98-03*. Wright-Patterson AFB, Ohio: Department of Electrical and Computer Engineering, Air Force Institute of Technology.
10. Zitzler, E., Deb, K., and Thiele, L. (2000) Comparison of multiobjective evolutionary algorithms: Empirical results. *Evolutionary Computation*, 8(2). 173–195.
11. Zitzler, E. and Thiele, L. (1998) Multiobjective optimization using evolutionary algorithms—A comparative case study. In Eiben, A. E., Bäck, T., Schoenauer, M., and Schwefel, H.-P., editors, *Parallel Problem Solving from Nature*, V, pages 292–301, Springer, Berlin, Germany.